

Programming Assignment 5

Varun Nagaraj

May 5th, 2019

1 Problem 1.a

Derive the restricted Boltzmann machine algorithm that you will implement, and explain your derivation. Implement the training and inference algorithms for RBM. Train RBMs with 20, 100 and 500 hidden nodes to generate MNIST images using the training data set. Generate MNIST images from the ones in the testing data set that have 20%, 50% and 80% pixels missing/removed. You are free to choose whether you want to use binary nodes or floating point nodes, but the derivation has to match the implementation.

Boltzmann Machines are a particular form of log-linear Markov Random Field i.e., for which the energy function is linear in its free parameters. To make them powerful enough to represent complicated distributions, we consider that some of the variables are hidden. By having more hidden variables, we can increase the modeling capacity of the Boltzmann Machine. Restricted Boltzmann Machines further restrict BMs to those without visible-visible and hidden-hidden connections.

The energy function $E(v,h)$ of an RBM is defined as:

$$E(v, h) = -b'v - c'h - h'Wv \quad (1)$$

Where,

W - Weights connecting the visible and hidden units;

b and c - offsets of the visible and hidden layers respectively;

This translates directly to the following free energy formula:

$$F(v) = -b'v - \sum_i \log \sum_{h_i} e^{h_i(c_i + W_i v)} \quad (2)$$

Because of the specific structure of RBMs, visible and hidden units are conditionally independent given one-another. Using this property, we can write:

$$p(h|v) = \prod_i p(h_i|v) \quad (3)$$

$$p(v|h) = \prod_j p(v_j|h) \quad (4)$$

A probabilistic version of the usual neuron activation function can be obtained by combining equation 1 with the below equation, where Z is called the partition function:

$$P(x) = \sum_h P(x, h) = \sum_h \frac{e^{-E(x, h)}}{Z} \quad (5)$$

Probabilistic version of neuron activation function using binary units where v_j and $h_i \in 0,1$ is as shown below:

$$P(h_i|v) = \text{sigm}(c_i + W_i v) \quad (6)$$

$$P(v_j|h) = \text{sigm}(b_j + W'_j h) \quad (7)$$

The free energy of an RBM with binary units further simplifies to:

$$F(v) = -b'v - \sum_i \log(1 + e^{(c_i + W_i v)}) \quad (8)$$

Samples used to estimate the negative phase gradient are referred to as negative particles, which are denoted as N . The gradient can then be written as:

$$-\frac{\partial \log p(x)}{\partial \theta} \approx \frac{\partial F(x)}{\partial \theta} - \frac{1}{|N|} \sum_{x' \in N} \frac{\partial F(x')}{\partial \theta} \quad (9)$$

Combining equations 8 and 9, we obtain the following log-likelihood gradients for an RBM with binary units as shown below:

$$\frac{\partial \log p(v)}{\partial W_{ij}} = E_v[p(h_i|v) * v_j] - v_j^i * \text{sigm}(W_i * v^i + c_j) \quad (10)$$

$$\frac{\partial \log p(v)}{\partial c_i} = E_v[p(h_i|v)] - \text{sigm}(W_i * v^i) \quad (11)$$

$$\frac{\partial \log p(v)}{\partial b_j} = E_v[p(v_j|h) * v_j] - v_j^i \quad (12)$$

2 Problem 1.b

Derive the variational autoencoder algorithm that you will implement, and explain your derivation. Implement the training and inference algorithms for VAE. Train VAE with 2, 8 and 16 code units to encode MNIST images using the training data set. The neural network will be 784 input -> 256 hidden -> 2/8/16 code -> 256 hidden -> 784 output. Then use the 2 code -> 256 hidden -> 784 output part of the trained network with 2 code units to generate images by varying each code unit from -3 to 3. You are free to choose the other parameters.

The idea of VAE is to identify $P(z)$ given $P(z|x)$. This can be done with the help of the KL divergence metric. Below is the proof which shows the encoder net, decoder net and the latent variable used in the auto-encoder.

Let's say we want to infer $P(z|x)$ given $Q(z|x)$. The KL divergence metric can be formulated

as follows:

$$D_{KL}[Q(z|X)||P(z|X)] = \Sigma_z Q(z|X) \log \frac{Q(z|X)}{P(z|X)} \quad (13)$$

$$D_{KL}[Q(z|X)||P(z|X)] = E[\log \frac{Q(z|X)}{P(z|X)}] \quad (14)$$

$$D_{KL}[Q(z|X)||P(z|X)] = E[\log Q(z|X) - \log P(z|X)] \quad (15)$$

By Bayes rule, we can transform the above equation to the one shown below

$$D_{KL}[Q(z|X)||P(z|X)] = E[\log Q(z|X) - \log \frac{P(X|z)P(z)}{P(X)}] \quad (16)$$

$$D_{KL}[Q(z|X)||P(z|X)] = E[\log Q(z|X) - (\log P(X|z) + \log P(z) - \log P(X))] \quad (17)$$

$$D_{KL}[Q(z|X)||P(z|X)] = E[\log Q(z|X) - \log P(X|z) - \log P(z) + \log P(X)] \quad (18)$$

$$D_{KL}[Q(z|X)||P(z|X)] - \log P(X) = E[\log Q(z|X) - \log P(X|z) - \log P(z)] \quad (19)$$

$$\log P(X) - D_{KL}[Q(z|X)||P(z|X)] = E[\log(P(X|z))] - E[\log Q(X|z) - \log P(z)] \quad (20)$$

$$\log P(X) - D_{KL}[Q(z|X)||P(z|X)] = E[\log(P(X|z))] - D_{KL}[Q(z|X)||P(z)] \quad (21)$$

Thus from equation 17 we understand that this is a structure of an Auto-encoder. In this the term $Q(z|X)$ can be treated as the encoder net, z is the encoded representation and $P(z|X)$ is the decoder net.

3 Code for part 1 derivation

```

1 # Created by Varun at 17/04/19
2 from __future__ import print_function
3 import numpy as np
4 import os
5 import random
6 import gzip, struct
7 import tensorflow as tf
8
9 def weight(shape, name='weights'):
10     return tf.Variable(tf.truncated_normal(shape, stddev=0.1), name=name)
11
12 def bias(shape, name='biases'):
13     return tf.Variable(tf.constant(0.1, shape=shape), name=name)
14
15 class RBM:
16     i = 0 # flipping index for computing pseudo likelihood
17
18     def __init__(self, n_visible=784, n_hidden=500, k=30, momentum=False):

```

```

19         self.n_visible = n_visible
20         self.n_hidden = n_hidden
21         self.k = k
22
23         self.lr = tf.placeholder(tf.float32)
24         if momentum:
25             self.momentum = tf.placeholder(tf.float32)
26         else:
27             self.momentum = 0.0
28         self.w = weight([n_visible, n_hidden], 'w')
29         self.hb = bias([n_hidden], 'hb')
30         self.vb = bias([n_visible], 'vb')
31
32         self.w_v = tf.Variable(tf.zeros([n_visible, n_hidden]), dtype=tf.float32)
33         self.hb_v = tf.Variable(tf.zeros([n_hidden]), dtype=tf.float32)
34         self.vb_v = tf.Variable(tf.zeros([n_visible]), dtype=tf.float32)
35
36     def propup(self, visible):
37         pre_sigmoid_activation = tf.matmul(visible, self.w) + self.hb
38         return tf.nn.sigmoid(pre_sigmoid_activation)
39
40     def prodown(self, hidden):
41         pre_sigmoid_activation = tf.matmul(hidden, tf.transpose(self.w)) + self.vb
42         return tf.nn.sigmoid(pre_sigmoid_activation)
43
44     def sample_h_given_v(self, v_sample):
45         h_props = self.propup(v_sample)
46         h_sample = tf.nn.relu(tf.sign(h_props - tf.random_uniform(tf.shape(h_props))))
47         return h_sample
48
49     def sample_v_given_h(self, h_sample):
50         v_props = self.prodown(h_sample)
51         v_sample = tf.nn.relu(tf.sign(v_props - tf.random_uniform(tf.shape(v_props))))
52         return v_sample
53
54     def CD_k(self, visibles):
55         # k steps gibbs sampling
56         v_samples = visibles
57         h_samples = self.sample_h_given_v(v_samples)
58         for i in range(self.k):
59             v_samples = self.sample_v_given_h(h_samples)
60             h_samples = self.sample_h_given_v(v_samples)
61
62         h0_props = self.propup(visibles)
63         w_positive_grad = tf.matmul(tf.transpose(visibles), h0_props)

```

```

64         w_negative_grad = tf.matmul(tf.transpose(v_samples), h_samples)
65         w_grad = (w_positive_grad - w_negative_grad) / tf.to_float(tf.shape(v
66         hb_grad = tf.reduce_mean(h0_props - h_samples, 0)
67         vb_grad = tf.reduce_mean(visibles - v_samples, 0)
68         return w_grad, hb_grad, vb_grad
69
70     def learn(self, visibles):
71         w_grad, hb_grad, vb_grad = self.CD_k(visibles)
72         # compute new velocities
73         new_w_v = self.momentum * self.w_v + self.lr * w_grad
74         new_hb_v = self.momentum * self.hb_v + self.lr * hb_grad
75         new_vb_v = self.momentum * self.vb_v + self.lr * vb_grad
76         # update parameters
77         update_w = tf.assign(self.w, self.w + new_w_v)
78         update_hb = tf.assign(self.hb, self.hb + new_hb_v)
79         update_vb = tf.assign(self.vb, self.vb + new_vb_v)
80         # update velocities
81         update_w_v = tf.assign(self.w_v, new_w_v)
82         update_hb_v = tf.assign(self.hb_v, new_hb_v)
83         update_vb_v = tf.assign(self.vb_v, new_vb_v)
84
85         return [update_w, update_hb, update_vb, update_w_v, update_hb_v, upda
86
87     def sampler(self, visibles, steps=5000):
88         v_samples = visibles
89         for step in range(steps):
90             v_samples = self.sample_v_given_h(self.sample_h_given_v(v_samples
91         return v_samples
92
93     def free_energy(self, visibles):
94         first_term = tf.matmul(visibles, tf.reshape(self.vb, [tf.shape(self.v
95         second_term = tf.reduce_sum(tf.log(1 + tf.exp(self.hb + tf.matmul(vis
96         return - first_term - second_term
97
98     def pseudo_likelihood(self, visibles):
99         x = tf.round(visibles)
100         x_fe = self.free_energy(x)
101         split0, split1, split2 = tf.split(x, [self.i, 1, tf.shape(x)[1] - sel
102         xi = tf.concat([split0, 1 - split1, split2], 1)
103         self.i = (self.i + 1) % self.n_visible
104         xi_fe = self.free_energy(xi)
105         return tf.reduce_mean(self.n_visible * tf.log(tf.nn.sigmoid(xi_fe - x
106
107 class DataSet:
108     batch_index = 0

```

109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153

```
def __init__(self, data_dir, batch_size = None, one_hot = False, seed = 0):
    self.data_dir = data_dir
    X, Y = self.read()
    shape = X.shape
    X = X.reshape([shape[0], shape[1] * shape[2]])
    self.X = X.astype(np.float)/255
    self.size = self.X.shape[0]
    if batch_size == None:
        self.batch_size = self.size
    else:
        self.batch_size = batch_size
    # abandon last few samples
    self.batch_num = int(self.size / self.batch_size)
    # shuffle samples
    np.random.seed(seed)
    np.random.shuffle(self.X)
    np.random.seed(seed)
    np.random.shuffle(Y)
    self.one_hot = one_hot
    if one_hot:
        y_vec = np.zeros((len(Y), 10), dtype=np.float)
        for i, label in enumerate(Y):
            y_vec[i, Y[i]] = 1.0
        self.Y = y_vec
    else:
        self.Y = Y

def read(self):
    with gzip.open(self.data_dir['Y']) as flbl:
        magic, num = struct.unpack(">II", flbl.read(8))
        label = np.fromstring(flbl.read(), dtype=np.int8)
    with gzip.open(self.data_dir['X'], 'rb') as fimg:
        magic, num, rows, cols = struct.unpack(">III", fimg.read(16))
        image = np.fromstring(fimg.read(), dtype=np.uint8).reshape(len(la
    return image, label

def next_batch(self):
    start = self.batch_index * self.batch_size
    end = (self.batch_index + 1) * self.batch_size
    self.batch_index = (self.batch_index + 1) % self.batch_num
    if self.one_hot:
        return self.X[start:end, :], self.Y[start:end, :]
    else:
        return self.X[start:end, :], self.Y[start:end]
```

```

154
155     def sample_batch(self):
156         index = random.randrange(self.batch_num)
157         start = index * self.batch_size
158         end = (index + 1) * self.batch_size
159         if self.one_hot:
160             return self.X[start:end, :], self.Y[start:end, :]
161         else:
162             return self.X[start:end, :], self.Y[start:end]
163
164     def random_removals(self, percentage):
165         index = random.randrange(self.batch_num)
166         start = index * self.batch_size
167         end = (index + 1) * self.batch_size
168         for i in self.X:
169             pixels = int(percentage*len(self.X[0]))
170             while(pixels > 0):
171                 rand = random.randint(0, 63)
172                 i[rand] = 0
173                 pixels -= 1
174                 if pixels <= 0:
175                     break
176             return self.X[start:end, :]
177
178 import scipy.misc
179 def save_images(images, size, path):
180     img = (images + 1.0) / 2.0
181     h, w = img.shape[1], img.shape[2]
182     merge_img = np.zeros((h * size[0], w * size[1]))
183     for idx, image in enumerate(images):
184         i = idx % size[1]
185         j = idx // size[1]
186         merge_img[j*h:j*h+h, i*w:i*w+w] = image
187     return scipy.misc.imsave(path, merge_img)
188
189
190 def train(train_data, test_data, epoches, percentage):
191     logs_dir = './logs'
192     samples_dir = './samples'
193
194     x = tf.placeholder(tf.float32, shape=[None, 784])
195     noise_x = test_data.random_removals(percentage)
196     hidden = 500
197     rbm = RBM(n_hidden=hidden)
198     step = rbm.learn(x)

```

```

199     sampler = rbm.sampler(x)
200     pl = rbm.pseudo_likelihood(x)
201
202     saver = tf.train.Saver()
203
204     with tf.Session() as sess:
205         init = tf.global_variables_initializer()
206         sess.run(init)
207         mean_cost = []
208         epoch = 1
209         for i in range(epochs * train_data.batch_num):
210             # draw samples
211             if i % 500 == 0:
212                 samples = sess.run(sampler, feed_dict = {x: noise_x})
213                 samples = samples.reshape([train_data.batch_size, 28, 28])
214                 save_images(samples, [8, 8], os.path.join(samples_dir, 'itera
215                     print('Saved samples. '))
216                 batch_x, _ = train_data.next_batch()
217                 sess.run(step, feed_dict = {x: batch_x, rbm.lr: 0.1})
218                 cost = sess.run(pl, feed_dict = {x: batch_x})
219                 mean_cost.append(cost)
220                 # save model
221                 if i is not 0 and train_data.batch_index is 0:
222                     checkpoint_path = os.path.join(logs_dir, 'model.ckpt')
223                     saver.save(sess, checkpoint_path, global_step = epoch + 1)
224                     print('Saved Model. ')
225                 # print pseudo likelihood
226                 if i is not 0 and train_data.batch_index is 0:
227                     print('Epoch %d Cost %g' % (epoch, np.mean(mean_cost)))
228                     mean_cost = []
229                     epoch += 1
230                 print('Test ')
231                 samples = sess.run(sampler, feed_dict = {x: noise_x})
232                 samples = samples.reshape([train_data.batch_size, 28, 28])
233                 save_images(samples, [8, 8], os.path.join(samples_dir, 'test{}-{}.png
234                 print('Saved samples. '))
235
236     train_dir = {
237         'X': './train-images-idx3-ubyte.gz',
238         'Y': './train-labels-idx1-ubyte.gz'
239     }
240
241     test_dir = {
242         'X': './t10k-images-idx3-ubyte.gz',
243         'Y': './t10k-labels-idx1-ubyte.gz'

```



```

244 }
245 train_data = DataSet(data_dir=train_dir, batch_size=64, one_hot=True)
246 test_data = DataSet(data_dir=test_dir, batch_size=64)
247 train(train_data, test_data, 3, 0.8)

```

4 Code for derivation 2

```

1 from __future__ import absolute_import
2 from __future__ import division
3 from __future__ import print_function
4
5 import argparse
6 import os
7
8 import matplotlib.pyplot as plt
9 import numpy as np
10 from keras import backend as K
11 from keras.datasets import mnist
12 from keras.layers import Lambda, Input, Dense
13 from keras.losses import mse, binary_crossentropy
14 from keras.models import Model
15 from keras.utils import plot_model
16
17
18 def sampling(args):
19     z_mean, z_log_var = args
20     batch = K.shape(z_mean)[0]
21     dim = K.int_shape(z_mean)[1]
22     # by default, random normal has mean=0 and std=1.0
23     epsilon = K.random_normal(shape=(batch, dim))
24     return z_mean + K.exp(0.5 * z_log_var) * epsilon
25
26
27 def plot_results(models,
28                 data,
29                 batch_size=128,
30                 model_name="vae_mnist"):
31     encoder, decoder = models
32     x_test, y_test = data
33     os.makedirs(model_name)
34
35     filename = os.path.join(model_name, "vae_mean.png")
36     z_mean, _, _ = encoder.predict(x_test,
37                                   batch_size=batch_size)
38     plt.figure(figsize=(12, 10))

```

```

39     plt.scatter(z_mean[:, 0], z_mean[:, 1], c=y_test)
40     plt.colorbar()
41     plt.xlabel("z[0]")
42     plt.ylabel("z[1]")
43     plt.savefig(filename)
44     plt.show()
45
46     filename = os.path.join(model_name, "digits_over_latent.png")
47     n = 30
48     digit_size = 28
49     figure = np.zeros((digit_size * n, digit_size * n))
50     grid_x = np.linspace(-3, 3, n)
51     grid_y = np.linspace(-3, 3, n)[::-1]
52
53     for i, yi in enumerate(grid_y):
54         for j, xi in enumerate(grid_x):
55             z_sample = np.array([[xi, yi]])
56             x_decoded = decoder.predict(z_sample)
57             digit = x_decoded[0].reshape(digit_size, digit_size)
58             figure[i * digit_size: (i + 1) * digit_size,
59                 j * digit_size: (j + 1) * digit_size] = digit
60
61     plt.figure(figsize=(10, 10))
62     start_range = digit_size // 2
63     end_range = n * digit_size + start_range + 1
64     pixel_range = np.arange(start_range, end_range, digit_size)
65     sample_range_x = np.round(grid_x, 1)
66     sample_range_y = np.round(grid_y, 1)
67     plt.xticks(pixel_range, sample_range_x)
68     plt.yticks(pixel_range, sample_range_y)
69     plt.xlabel("z[0]")
70     plt.ylabel("z[1]")
71     plt.imshow(figure, cmap='Greys_r')
72     plt.savefig(filename)
73     plt.show()
74
75
76 (x_train, y_train), (x_test, y_test) = mnist.load_data()
77 image_size = x_train.shape[1]
78 original_dim = image_size * image_size
79 x_train = np.reshape(x_train, [-1, original_dim])
80 x_test = np.reshape(x_test, [-1, original_dim])
81 x_train = x_train.astype('float32') / 255
82 x_test = x_test.astype('float32') / 255
83 input_shape = (original_dim,)

```

```

84 intermediate_dim = 256
85 batch_size = 128
86 latent_dim = 2
87 epochs = 50
88
89 inputs = Input(shape=input_shape, name='encoder_input')
90 x = Dense(intermediate_dim, activation='relu')(inputs)
91 z_mean = Dense(latent_dim, name='z_mean')(x)
92 z_log_var = Dense(latent_dim, name='z_log_var')(x)
93
94 z = Lambda(sampling, output_shape=(latent_dim,), name='z')([z_mean, z_log_var])
95 encoder = Model(inputs, [z_mean, z_log_var, z], name='encoder')
96 encoder.summary()
97 plot_model(encoder, to_file='vae_mlp_encoder.png', show_shapes=True)
98 latent_inputs = Input(shape=(latent_dim,), name='z_sampling')
99 x = Dense(intermediate_dim, activation='relu')(latent_inputs)
100 outputs = Dense(original_dim, activation='sigmoid')(x)
101 decoder = Model(latent_inputs, outputs, name='decoder')
102 decoder.summary()
103 plot_model(decoder, to_file='vae_mlp_decoder.png', show_shapes=True)
104 outputs = decoder(encoder(inputs)[2])
105 vae = Model(inputs, outputs, name='vae_mlp')
106
107 if __name__ == '__main__':
108     parser = argparse.ArgumentParser()
109     help_ = "Load h5 model trained weights"
110     parser.add_argument("-w", "--weights", help=help_)
111     help_ = "Use mse loss instead of binary cross entropy (default)"
112     parser.add_argument("-m",
113                         "--mse",
114                         help=help_, action='store_true')
115     args = parser.parse_args()
116     models = (encoder, decoder)
117     data = (x_test, y_test)
118     if args.mse:
119         reconstruction_loss = mse(inputs, outputs)
120     else:
121         reconstruction_loss = binary_crossentropy(inputs,
122                                                  outputs)
123
124     reconstruction_loss *= original_dim
125     kl_loss = 1 + z_log_var - K.square(z_mean) - K.exp(z_log_var)
126     kl_loss = K.sum(kl_loss, axis=-1)
127     kl_loss *= -0.5
128     vae_loss = K.mean(reconstruction_loss + kl_loss)

```


```

129     vae.add_loss(vae_loss)
130     vae.compile(optimizer='adam')
131     vae.summary()
132     plot_model(vae,
133                to_file='vae_mlp.png',
134                show_shapes=True)
135
136     if args.weights:
137         vae.load_weights(args.weights)
138     else:
139         vae.fit(x_train,
140                epochs=epochs,
141                batch_size=batch_size,
142                validation_data=(x_test, None))
143         vae.save_weights('vae_mlp_mnist.h5')
144
145     plot_results(models,
146                 data,
147                 batch_size=batch_size,
148                 model_name="vae_mlp")


```

5 Code Output images


5.1 Part-1




- 5.1.1 The first row is with 20 hidden nodes, second is with 100 and third is with 500 hidden nodes. The first column shows the output when 20 percent of the pixels are removed, similarly column 2- 50 percent and column 3-80 percent



5.2 Part-2





6 References

<https://jaan.io/what-is-variational-autoencoder-vae-tutorial/>
<https://medium.com/datatype/restricted-boltzmann-machine-a-complete-analysis-part-1-introduction-model-formulation-1a4404873b3>
<https://towardsdatascience.com/teaching-a-variational-autoencoder-vae-to-draw-mnist-characters-978675c95776>
<https://github.com/FelixMohr/Deep-learning-with-Python/blob/master/VAE.ipynb>
<https://machinelearningmastery.com/tutorial-first-neural-network-python-keras/>
https://oeis.org/wiki/List_of_LaTeX_mathematical_symbols#Relation_operators
<https://medium.com/datadriveninvestor/image-processing-for-mnist-using-keras-f9a1021f6ef0>
https://github.com/keras-team/keras/blob/master/examples/variational_autoencoder.py